

Red-Black Trees *Lecture*

Instructor: Dr. G P Gupta

Red-black trees: Overview

- ♦ Red-black trees are a variation of binary search trees to ensure that the tree is **balanced**.
 - » Height is $O(\lg n)$, where n is the number of nodes.
- ♦ Operations take $O(\lg n)$ time in the **worst case**.

Red-black Tree

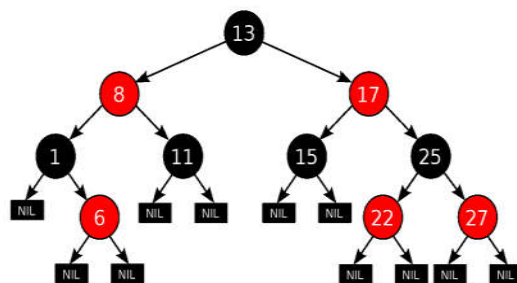
- ♦ A **red-black tree** is a **binary search tree with one extra bit of storage** per node: its color, which can be either **RED** or **BLACK**.
- ♦ All other attributes of BSTs are inherited:
 - » *key*, *left*, *right*, and *p*.
- ♦ All empty trees (leaves) are colored black.
 - » We use a single sentinel, *nil*, for all the leaves of red-black tree T , with $\text{color}[\text{nil}] = \text{black}$.
 - » The root's parent is also $\text{nil}[T]$.

Red-black Properties

A red-black tree is a binary tree that satisfies the following red-black properties:

1. Every node is either **red** or **black**.
2. The **root** is **black**.
3. Every **leaf (nil)** is **black**.
4. If a node is **red**, then both its children are **black**.
5. For each node, All simple paths from any node x to a descendant leaf have the same number of black nodes = **black-height(x)**.

Red-black Tree – Example

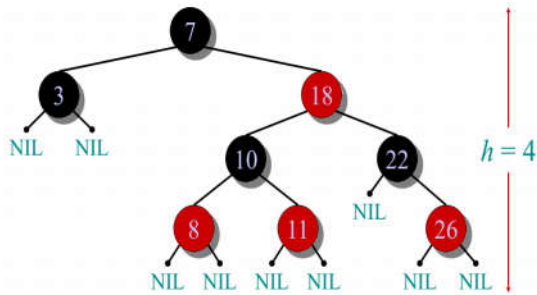


An example of a red-black tree

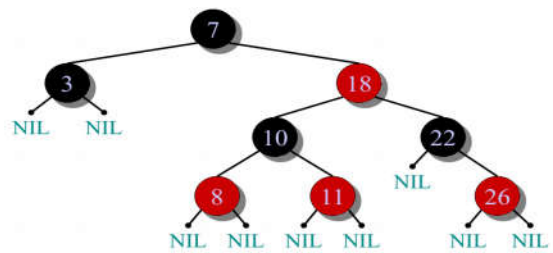
Height of a Red-black Tree

- ♦ **Height of a node:**
 - » $h(x)$ = number of edges in a longest path to a leaf.
- ♦ **Black-height of a node x , $bh(x)$:**
 - » $bh(x)$ = number of black nodes (including $\text{nil}[T]$) on the path from x to leaf, not counting x .
- ♦ **Black-height of a red-black tree is the black-height of its root.**
 - » By Property 5, **black height is well defined**.

Example of a red-black tree

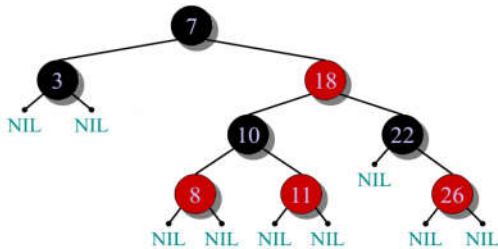


Example of a red-black tree



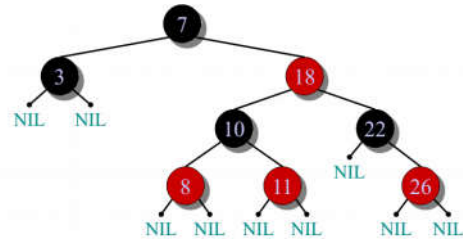
1. Every node is either red or black.

Example of a red-black tree



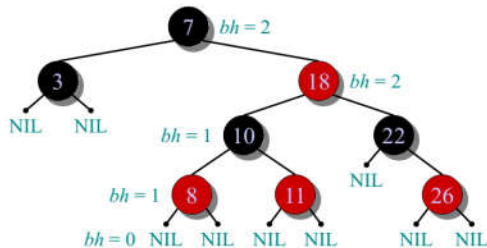
2. The root and leaves (NIL's) are black.

Example of a red-black tree



3. If a node is red, then its parent is black.

Example of a red-black tree



4. All simple paths from any node x to a descendant leaf have the same number of black nodes = **black-height**(x).

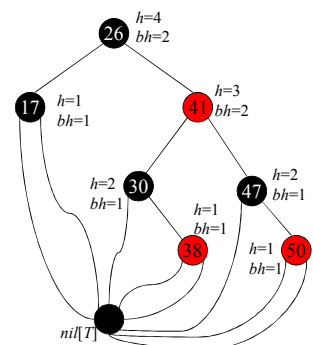
Height of a Red-black Tree

- Example:

- **Height of a node:**
 $h(x)$ = # of edges in a longest path to a leaf.

- **Black-height of a node**
 $bh(x)$ = # of black nodes on path from x to leaf, not counting x .

- How are they related?
» $bh(x) \leq h(x) \leq 2 bh(x)$



red-black tree

- ♦ **Lemma:** *The subtree rooted at any node x has $\geq 2^{bh(x)} - 1$ internal nodes.*
- ♦ **Proof:** By induction on height of x .
 - » **Base Case:** Height $h(x) = 0 \Rightarrow x$ is a leaf $\Rightarrow bh(x) = 0$.
Subtree has $2^0 - 1 = 0$ nodes. ✓
 - » **Induction Step:** Height $h(x) = h > 0$ and $bh(x) = b$.
 - Each child of x has height $h - 1$ and black-height either b (child is **red**) or $b - 1$ (child is **black**).
 - **By Induction hypothesis**, each child has $\geq 2^{bh(x)-1} - 1$ internal nodes.
 - **Subtree rooted at x has** $\geq 2(2^{bh(x)-1} - 1) + 1 = 2^{bh(x)} - 1$ internal nodes. (**The + 1 is for x itself.**)

Height of a red-black tree

Theorem. A red-black tree with n keys has height $h \leq 2 \lg(n + 1)$.

Proof:

- **Lemma:** The subtree rooted at any node x has $\geq 2^{bh(x)} - 1$ internal nodes.
- By the above lemma, $n \geq 2^{bh} - 1$,
and since $bh \geq h/2$, we have $n \geq 2^{h/2} - 1$.
 $\Rightarrow h \leq 2 \lg(n + 1)$.

Lemma “RB Height”

Lemma 1: Consider a node x in an RB tree: The longest descending path from x to a leaf has length $h(x)$, which is at most twice the length of the shortest descending path from x to a leaf.

Proof:

black nodes on any path from $x = bh(x)$ (prop 5)
 \leq # nodes on shortest path from x , $s(x)$. (prop 1)
 But, there are no consecutive red (prop 4),
 and we end with black (prop 3), so $h(x) \leq 2 bh(x)$.
Thus, $h(x) \leq 2 s(x)$.

Operations on RB Trees

- ♦ **Insertion and Deletion**
- ♦ **The queries :** *SEARCH, MIN, MAX, SUCCESSOR, and PREDECESSOR*
 - all run in $O(\lg n)$ time on a red-black tree with n nodes.
- ♦ All operations can be performed in **$O(\lg n)$** time.

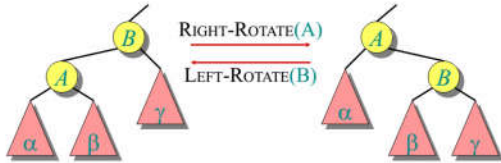
Operations on RB Trees

- ♦ Insertion and Deletion *are not straightforward.* Why?

Operations on RB Trees

- ♦ The operations INSERT and DELETE cause modifications to the red-black tree:
 - **color changes,**
 - **restructuring the links of the tree: “rotations”.**

Rotations



Rotations maintain the inorder ordering of keys:

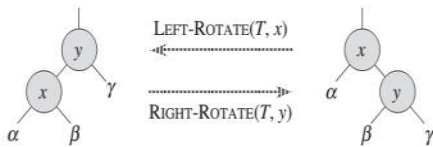
• $a \in \alpha, b \in \beta, c \in \gamma \Rightarrow a \leq A \leq B \leq c$.

A rotation can be performed in $O(1)$ time.

Rotation

- The pseudo-code for Left-Rotate assumes that
 - » $right[x] \neq nil[T]$, and
 - » root's parent is $nil[T]$.
- Left Rotation on x , makes x the left child of y , and the left subtree of y into the right subtree of x .
- Pseudocode for Right-Rotate is symmetric: exchange *left* and *right* everywhere.
- Time:** $O(1)$ for both Left-Rotate and Right-Rotate, since a constant number of pointers are modified.

LEFT-ROTATE

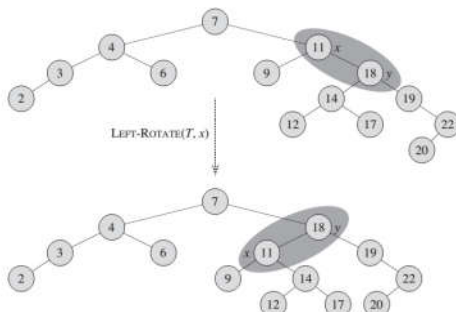


LEFT-ROTATE

```

LEFT-ROTATE( $T, x$ )
1   $y = x.right$            // set y
2   $x.right = y.left$        // turn y's left subtree into x's right subtree
3  if  $y.left \neq T.nil$ 
4       $y.left.p = x$ 
5   $y.p = x.p$              // link x's parent to y
6  if  $x.p == T.nil$ 
7       $T.root = y$ 
8  elseif  $x == x.p.left$ 
9       $x.p.left = y$ 
10 else  $x.p.right = y$ 
11  $y.left = x$            // put x on y's left
12  $x.p = y$ 
    
```

LEFT-ROTATE



Insertion in RB Trees

- Insertion **must preserve** all red-black properties.
- Should an inserted node be colored **Red**? Black?
- Basic steps:**
 - » Use **Tree-Insert from BST** (slightly modified) to insert a node z into T .
 - Procedure **RB-Insert(T, z)**.
 - » Color the node z **red**.
 - » Fix the modified tree **by re-coloring nodes** and performing **rotation** to preserve RB tree property.
 - Procedure **RB-Insert-Fixup**.

RB-INSERT(T, z)

```

1   $y = T.nil$ 
2   $x = T.root$ 
3  while  $x \neq T.nil$ 
4       $y = x$ 
5      if  $z.key < x.key$ 
6           $x = x.left$ 
7      else  $x = x.right$ 
8   $z.p = y$ 
9  if  $y == T.nil$ 
10      $T.root = z$ 
11  elseif  $z.key < y.key$ 
12      $y.left = z$ 
13  else  $y.right = z$ 
14   $z.left = T.nil$ 
15   $z.right = T.nil$ 
16   $z.color = RED$ 
17  RB-INSERT-FIXUP( $T, z$ )

```

RB-INSERT-FIXUP(T, z)

```

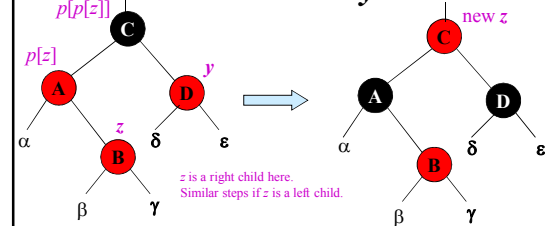
1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == RED$ 
5               $z.p.color = BLACK$  // case 1
6               $y.color = BLACK$  // case 1
7               $z.p.p.color = RED$  // case 1
8               $z = z.p.p$  // case 1
9          else if  $z == z.p.p.right$ 
10              $z = z.p$  // case 2
11             LEFT-ROTATE( $T, z$ ) // case 2
12              $z.p.color = BLACK$  // case 3
13              $z.p.p.color = RED$  // case 3
14             RIGHT-ROTATE( $T, z.p.p$ ) // case 3
15         else (same as then clause
16             with "right" and "left" exchanged)
17      $T.root.color = BLACK$ 

```

Insertion into a red-black tree

- There are three Case:

Case 1 – uncle y is red



- $p[p[z]]$ (z 's grandparent) must be black, since z and $p[z]$ are both red and there are no other violations of property 4.
- Make $p[z]$ and y black \Rightarrow now z and $p[z]$ are not both red. But property 5 might now be violated.
- Make $p[p[z]]$ red \Rightarrow restores property 5.
- The next iteration has $p[p[z]]$ as the new z (i.e., z moves up 2 levels).

Insertion into a red-black tree

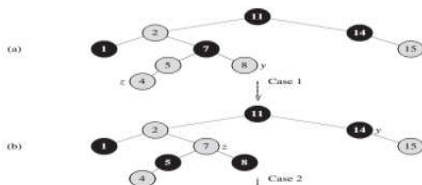
- Case 1

RB-INSERT-FIXUP(T, z)

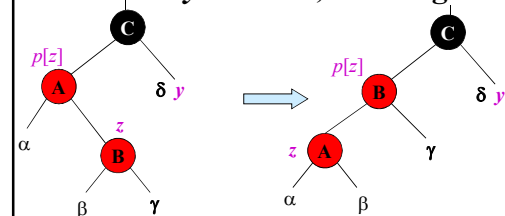
```

1  while  $z.p.color == RED$ 
2      if  $z.p == z.p.p.left$ 
3           $y = z.p.p.right$ 
4          if  $y.color == RED$ 
5               $z.p.color = BLACK$  // case 1
6               $y.color = BLACK$  // case 1
7               $z.p.p.color = RED$  // case 1
8               $z = z.p.p$  // case 1

```



Case 2 – y is black, z is a right child



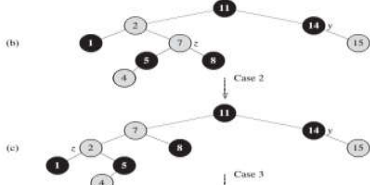
- Left rotate around $p[z]$** , $p[z]$ and z switch roles \Rightarrow now z is a left child, and both z and $p[z]$ are red.
- Takes us immediately to case 3.

Insertion into a red-black tree

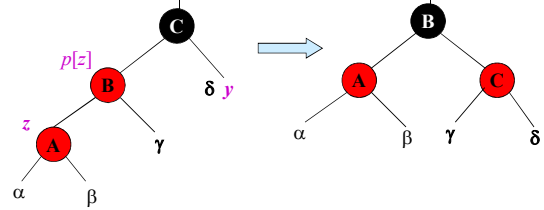
Case 2

```

RB-INSERT-FIXUP(T, z)
1 while z.p.color == RED
2   if z.p == z.p.p.left
3     y = z.p.p.right
4     if y.color == RED
5       z.p.color = BLACK // case 1
6       y.color = BLACK // case 1
7       z.p.p.color = RED // case 1
8       z = z.p.p
9     else if z == z.p.right
10      z = z.p // case 2
11      LEFT-ROTATE(T, z) // case 2
    
```



Case 3 – y is black, z is a left child



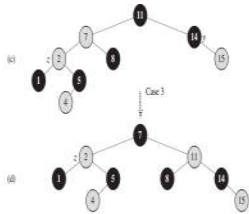
- Make $p[z]$ black and $p[p[z]]$ red.
- Then **right rotate on $p[p[z]]$** . Ensures property 4 is maintained.
- No longer have 2 reds in a row.
- $p[z]$ is now black \Rightarrow no more iterations.

Insertion into a red-black tree

Case 3

```

RB-INSERT-FIXUP(T, z)
1 while z.p.color == RED
2   if z.p == z.p.p.left
3     y = z.p.p.right
4     if y.color == RED
5       z.p.color = BLACK // case 1
6       y.color = BLACK // case 1
7       z.p.p.color = RED // case 1
8       z = z.p.p
9     else if z == z.p.right
10      z = z.p // case 2
11      LEFT-ROTATE(T, z) // case 2
12   z.p.color = BLACK // case 3
13   z.p.p.color = RED // case 3
14   RIGHT-ROTATE(T, z.p.p) // case 3
15   else (same as then clause with "right" and "left" exchanged)
16   T.root.color = BLACK
    
```



Algorithm Analysis

- $O(\lg n)$ time to get through RB-Insert up to the call of RB-Insert-Fixup.
- Within RB-Insert-Fixup:**
 - Each iteration takes $O(1)$ time.
 - Each iteration but the last **moves z up 2 levels**.
 - $O(\lg n)$ levels $\Rightarrow O(\lg n)$ time.
 - Thus, **insertion in a red-black tree takes $O(\lg n)$ time**.
 - Note: **there are at most 2 rotations overall**.

Exercise

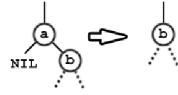
- Show the red-black trees that result after successively inserting the keys 41; 38; 31; 12; 19; 8 into an initially empty red-black tree.

red-black tree :Deletion

- Deleting a node from a red-black tree is a bit more complicated than inserting a node.

red-black tree :Deletion

- Use usual BST deletion algorithm.
- Let deleting node **a** (disregard colors, fix later)
- **Case 1: a has no left child**

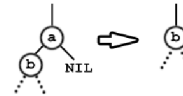


- Remove a and put its right child (b) instead
- **Note:**
 - if the red rule is now broken b and its new father (originally a's father), we can color node b in black; keeping the black height balance, since a was definitely black (as its father is red)

red-black tree :Deletion cont..

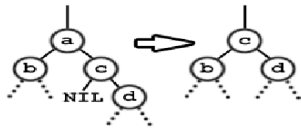
- **Case 2: a has no right child:**
- Remove a and put its left child b in its place

Note: same as case 1



red-black tree :Deletion cont..

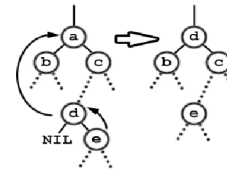
- **Case 3: a has two children, a's successor (c) is its right child:**
 - » Remove a and put its successor (c) in its place
 - » Make a's left child (b) the successor's (c) left child



- **Note:**
 - successor node always has no left child
 - Moving the successor node, we color it in a's color.
 - If the successor was black, the child that replaced it (d) is colored in "extra" black, making it red-black or black-black. This is fixed in the correction.

red-black tree :Deletion cont..

- **Case 4: a has two children, a's successor (d) is not its child:**
 - » Put the successor's (d) left child (e) instead of it
 - » Remove a and put its successor (d) instead of it,
 - » making a's children (b, c) its new children



Notes: same as case 3

red-black tree :Deletion

- When we want to delete node z and **z has fewer than two children**, then z is removed from the tree, and then assign z to the y.
- When **z has two children**, then y should be z's successor, and y moves into z's position in the tree.
 - » remember y's color before it is removed from or moved within the tree, and
 - » keep track of the node x that moves into y's original position in the tree, because node x might also cause violations of the red-black properties.
- After deleting node z, RB-DELETE calls an auxiliary procedure **RB-DELETE-FIXUP**, which changes colors and performs rotations to restore the red-black properties.

red-black tree :Deletion

```

RB-DELETE(T, z)
1  y = z
2  y-original-color = y.color
3  if z.left == T.nil
4      x = z.right
5      RB-TRANSPLANT(T, z, z.right)
6  elseif z.right == T.nil
7      x = z.left
8      RB-TRANSPLANT(T, z, z.left)
9  else y = TREE-MINIMUM(z.right)
10 y-original-color = y.color
11 x = y.right
12 if y.p == z
13     x.p = y
14 else RB-TRANSPLANT(T, y, y.right)
15     y.right = z.right
16     y.right.p = y
17 RB-TRANSPLANT(T, z, y)
18 y.left = z.left
19 y.left.p = y
20 y.color = z.color
21 if y-original-color == BLACK
22     RB-DELETE-FIXUP(T, x)
    
```

Correction after Deletion in RB-Tree

RB-DELETE-FIXUP

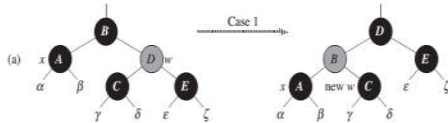
```

RB-DELETE-FIXUP(T, x)
1  while x ≠ T.root and x.color == BLACK
2    if x == x.p.left
3      w = x.p.right
4      if w.color == RED
5        w.color = BLACK
6        x.p.color = RED
7        LEFT-ROTATE(T, x.p)
8        w = x.p.right
9      if w.left.color == BLACK and w.right.color == BLACK
10       w.color = RED
11       x = x.p
12     else if w.right.color == BLACK
13       w.left.color = BLACK
14       w.color = RED
15       RIGHT-ROTATE(T, w)
16       w = x.p.right
17       w.color = x.p.color
18       x.p.color = BLACK
19       w.right.color = BLACK
20       LEFT-ROTATE(T, x.p)
21       x = T.root
22     else (same as then clause with "right" and "left" exchanged)
23   x.color = BLACK
    
```

RB-DELETE-FIXUP : case-1

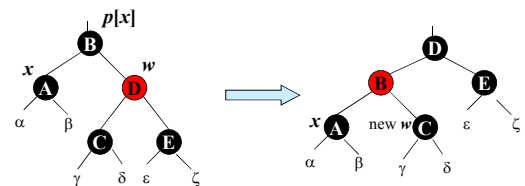
```

RB-DELETE-FIXUP(T, x)
1  while x ≠ T.root and x.color == BLACK
2    if x == x.p.left
3      w = x.p.right
4      if w.color == RED
5        w.color = BLACK
6        x.p.color = RED
7        LEFT-ROTATE(T, x.p)
8        w = x.p.right
    
```



Case 1: x's sibling w is red

Case 1 – w is red



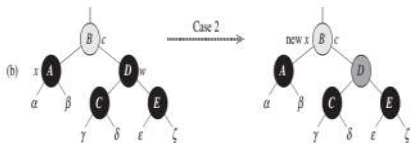
- w must have black children.
- Make w black and p[x] red (because w is red p[x] couldn't have been red).
- Then left rotate on p[x].
- New sibling of x was a child of w before rotation \Rightarrow must be black.
- Go immediately to case 2, 3, or 4.

RB-DELETE-FIXUP : case-2

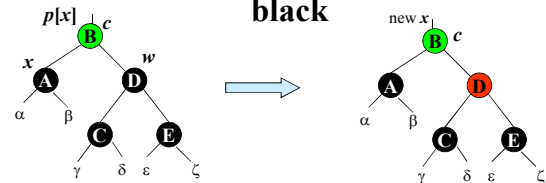
Case 2: x's sibling w is black, and both of w's children are black

```

9  if w.left.color == BLACK and w.right.color == BLACK
10   w.color = RED
11   x = x.p
    
```



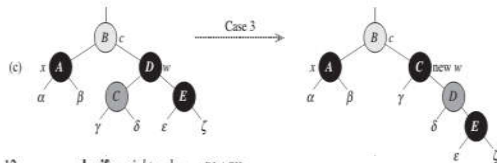
Case 2 – w is black, both w's children are black



- Take 1 black off x (\Rightarrow singly black) and off w (\Rightarrow red).
- Move that black to p[x].
- Do the next iteration with p[x] as the new x.
- If entered this case from case 1, then p[x] was red \Rightarrow new x is red & black \Rightarrow color attribute of new x is RED \Rightarrow loop terminates. Then new x is made black in the last line.

RB-DELETE-FIXUP : case-3

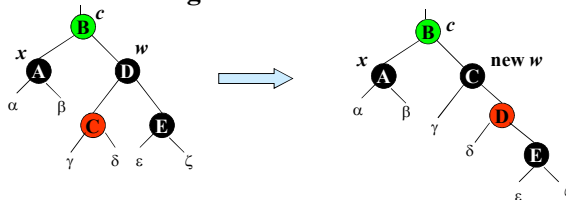
- Case 3: x 's sibling w is black, w 's left child is red, and w 's right child is black



```

12     else if w.right.color == BLACK
13         w.left.color = BLACK           // case 3
14         w.color = RED                  // case 3
15         RIGHT-ROTATE(T, w)             // case 3
16         w = x.p.right                  // case 3
    
```

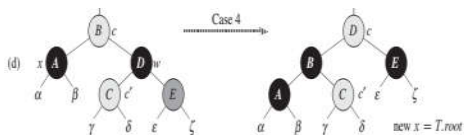
Case 3 – w is black, w 's left child is red, w 's right child is black



- Make w red and w 's left child black.
- Then right rotate on w .
- New sibling w of x is black with a red right child \Rightarrow case 4.

RB-DELETE-FIXUP : case-4

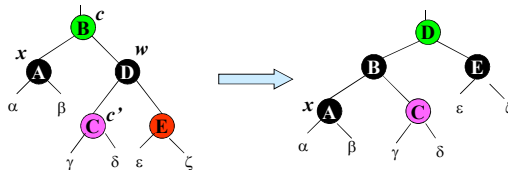
- Case 4: x 's sibling w is black, and w 's right child is red



```

17         w.color = x.p.color           // case 4
18         x.p.color = BLACK              // case 4
19         w.right.color = BLACK          // case 4
20         LEFT-ROTATE(T, x.p)           // case 4
21         x = T.root                    // case 4
22     else (same as then clause with "right" and "left" exchanged)
23         x.color = BLACK
    
```

Case 4 – w is black, w 's right child is red



- Make w be $p[x]$'s color (c).
- Make $p[x]$ black and w 's right child black.
- Then left rotate on $p[x]$.
- Remove the extra black on x ($\Rightarrow x$ is now singly black) without violating any red-black properties.
- All done. Setting x to root causes the loop to terminate.

Analysis

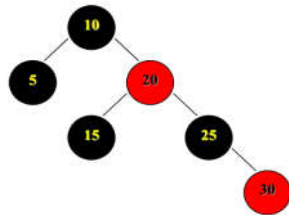
- $O(\lg n)$ time to get through RB-Delete up to the call of RB-Delete-Fixup.
- Within RB-Delete-Fixup:
 - Case 2 is the only case in which more iterations occur.
 - x moves up 1 level.
 - Hence, $O(\lg n)$ iterations.
 - Each of cases 1, 3, and 4 has 1 rotation $\Rightarrow \leq 3$ rotations in all.
 - Hence, $O(\lg n)$ time.

Exercises-1

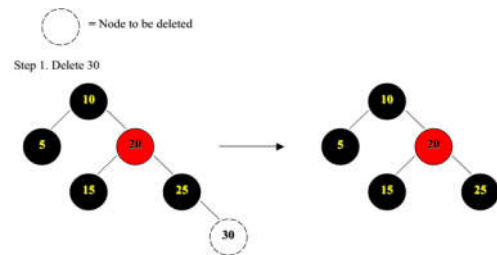
- Suppose that a node x is inserted into a red-black tree with RB-INSERT and then immediately deleted with RB-DELETE. Is the resulting red-black tree the same as the initial red-black tree? Justify your answer.

Exercises-2

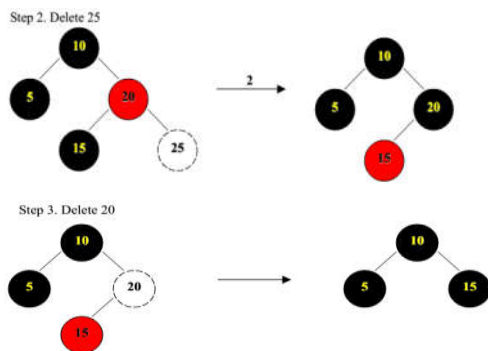
- ♦ Show the red-black tree that results from successively deleting the keys 30, 25, 20, 15, 10, and 5 from following RB-Tree.



Exercises-2: solution



Exercises-2: solution



Exercises-3: solution

